

System memory - physical memory are some hardware

Multiplexing

Protection: Protection

Controlled overlap

Translation:

Virtual memory \rightarrow physical memory

$$EAT = \text{Hit Time} + \text{Miss Rate} \times \text{Miss Penalty}$$

Addressing Approaches

Unpronouncing only one program in the memory once

Not efficient (High overhead of CS in multiple program)

Not powerful: Cannot run multiple programs at a time

Whispering allows multiple programs using memory

More complex in OS (需要多块, isolated memory & translation)

地址 (fixed-sized partition) 地址 (dynamic partition)

虚拟 Memory: Base & Bound • Segmentation

Paging + Segmentation

Virtual memory works? CPU fetch instruction

The memory returns the instruction

Processor will do not need all memory within allocated chunk

do not need all memory within allocated chunk

Segmentation

将内存分割成大小不同的块, 可以根据访问模式重新组合

容易共享, 容易保护, 减少碎片

但: 复杂, 增加开销

增加寻址时间

增加访存时间

减少碎片

提高内存利用率

减少碎片

提高内存利用率

In/Out or load/store

Pro: Simple hardware, easy to program

Con: Consumes processor cycles proportional to data size

Direct Memory Access Give controller access to memory

bus, ask it to transfer data blocks to/ from memory, free of disks into files, Directories, etc.

bio, tell device driver to transfer disk data to buffer

bio tell device driver to transfer disk data to buffer

at address X • Device driver tells disk controller to transfer

at address X • Device driver to buffer at address X • Disk controller send each byte to Reliability: keep files despite crashes, failure, etc.

Layout contiguous, linked,inode, In/Location

Continuous Allocation

at addresses X, increasing memory address and decreasing C until C=0

Pro: Locate files easily, File deletion easily

Con: external fragmentation

Linked Allocation

Pro: reduce fragmentation, flexible

Con: poor random access performance

FAT

Root sector FSINFO FAT1 FAT2 Directory Root File&Data

• Read root directory, retrieve first block number

• Read FAT determine the location of next block

• Read FAT determine the location of next block

• The process stops until FAT says the last block

FAT16 = 2¹⁶ FAT32 = 2³² clusters

FAT12 = 2¹²

bytes

Largest size of a FAT32 file = 4G - 1

File system size = 2³² * block size

FAT32 File system entry

FAT series - directory entry

stores the start cluster number

stores the file size

FAT series - reading a file

• Read the content from first cluster

• Look for the next cluster until reach the last

FAT series - writing a file

• Locate the last cluster

• Start writing to the non-full cluster

• Allocate the next cluster through FSINFO

• Update the FAT and FSINFO

• When writing finishes, update the file size

File system

- deleting a file

• Delete all the blocks involved.

• Update FAT and FSINFO

• Change first byte of the direct entry to - (0xE5)

FAT scheme

• Transform blocks into files and directories

• Optimize for access and usage pattern

• Maximize sequential access, allow efficient random access

Link list approach

Concurrently: cameras, USB drives, SD cards

Natively used: memory, I/O device service time

Controlling factor to latency

Software Both: Hardware controller, I/O device service time

Instruction - I/O: 1. I/O 2. I/O instruction

- I/O instruction

Cache Register

- Cache memory, need more time to access memory

Content Switch

Content of top-level segment registers

Pointer to top-level table

Pointer to top-level table

Cache: A repository for copies that can be accessed more quickly than the origin

Access Time

Latency: Time taken to access data items closer to processor

Access Time

Latency: Time taken to access data items closer to processor

Access Time

Latency: Time taken to access data items closer to processor

Access Time

Block size \times Black size ($\frac{\text{Addr Length}}{\text{Offset}}$)

Max file size ! = FS size

FS Layout			
File	Free	Block	inode
Block P	Free	Block	inode Table
Block T	GDT	Bitmap	Blocks

by Ext. having group:
Performance: spatial locality
Reliability: superblock and GDT are replicated in each and never return to the origin code. Will replace program code, memory local v. global v. dynamical memory), free values { Multi programming A \rightarrow B } switching thread across blocks require changes to lock (int process) if interested [process] = FALSE; };

• **System Call • interrupt • Trap or Exception**

fork(1): It creates the child process by cloning from the parent process, including all user-space data. By [CPU] real, however, do not clone return value of fork(), PID, parent process, running time, file locks.

area: The process is changing the code that is executing. Will replace program code, memory local v. global v. dynamical memory), free values { Multi programming A \rightarrow B } switching thread across blocks require changes to lock (int process) if interested [process] = TRUE; turn = other; while turn = other wait unlock (int process) if interested [process] = FALSE; };

MultiProcessing = Multiple CPU
Multiprogramming = Multiple Job / Process
Multithreading = Multiple thread per process
Memory and ID Address table

producer-consumer
A Low priority process L is inside the critical region, but a high priority process H gets the CPU and want to enter the critical region. But H cannot lock (since L is not awake). wait (mutex); wait (mutex); insert item;

producer-consumer
Producer mutex = 1; avail = N; full = 0;
Consumer mutex = 1; int item = producer item; while (true) item = producer int item; while (true) item = producer item; insert (item); wait (full); post (full); ?;

EFS: void consumer (int item; while (true) item = producer item; wait (full); post (full); ?;

ES: void producer (int item; insert (item); wait (avail); post (full); ?;

Dining Philosopher (LEFT ((i+n-1)%N Right (i+1)%N). Right (i+1)%N. Left state[N]; semaphore mutex = 1. PL/N>=0
NF: void philosopher (int i) {think(); take(i); eat(); put(i);?; if (i==n-1) i=i+1; else i=i+1; };

SE: void take (int i) { wait (&mutex); state[i] = hungry; captain[i] = post (&mutex); wait (&mutex); state[i] = Think; captain[i] = post (&mutex); wait (&mutex); state[i] = Hungry; captain[LEFT] = captain[RIGHT]; post (&mutex);?; if (i==1) i=i+1; else i=i+1; };

HF: void captain (int i) { if (state[i] == E; post (&mutex);?; if (state[i] != E) i=i+1; else i=i+1; };

State LF: ! = E; if (state[i] == E; post (&mutex);?; if (state[i] != E;

Wait & post
when SIGCHLD comes, the corresponding signal handling routine is registered in the parent with SIGCHLD. When SIGCHLD is issued, the parent will receive the signal. If it could be achieved, the problem of race exclusion is a requirement of race condition. If it could be avoided, the problem of race condition is solved.

critical section: A code segment that access shared objects simultaneously and can be designed for accessing more than one shared object.

File system calls -> **Entry and Exit Implementation**

fs
1. clean up the exit process's resource.
2. clean up the exit process's resource.
3. notify the parent with SIGCHLD when SIGCHLD is issued. The corresponding signal handling routine is invoked (remove the signal, destroy the child process condition variable). If it could be avoided, the problem of race exclusion is solved.

fs (entry)
1. deregister the signal handling routine in the kernel space.
2. for the parent 4. return the PID of child for the parent 4. return the PID of child for the parent 4. return the PID of child for the parent 4. return the PID of child for the parent 4. return the PID of child

fs (exit)
1. The process terminates abnormally.
2. The process terminates normally.
3. The process terminates normally.

When making a system call (fs) switches from user mode to kernel mode.

when making a system call (fs) switches from user mode to kernel mode.

Time

User Time

CPU

idle

IO

SWI

SWI

SWI

SWI

3 types of Mode transfer: Unprogrammed

• System Call • interrupt • Trap or Exception

fork(1): It creates the child process by cloning from the parent process, including all user-space data. By [CPU] real, however, do not clone return value of fork(), PID, parent process, running time, file locks.

area: The process is changing the code that is executing. Will replace program code, memory local v. global v. dynamical memory), free values { Multi programming A \rightarrow B } switching thread across blocks require changes to lock (int process) if interested [process] = TRUE; turn = other; while turn = other wait unlock (int process) if interested [process] = FALSE; };

producer-consumer
A Low priority process L is inside the critical region, but a high priority process H gets the CPU and want to enter the critical region. But H cannot lock (since L is not awake). wait (mutex); wait (mutex); insert item;

producer-consumer
Producer mutex = 1; avail = N; full = 0;
Consumer mutex = 1; int item = producer item; while (true) item = producer int item; while (true) item = producer item; insert (item); wait (full); post (full); ?;

EFS: void consumer (int item; while (true) item = producer item; wait (full); post (full); ?;

ES: void producer (int item; insert (item); wait (avail); post (full); ?;

Dining Philosopher (LEFT ((i+n-1)%N Right (i+1)%N). Right (i+1)%N. Left state[N]; semaphore mutex = 1. PL/N>=0
NF: void philosopher (int i) {think(); take(i); eat(); put(i);?; if (i==n-1) i=i+1; else i=i+1; };

SE: void take (int i) { wait (&mutex); state[i] = hungry; captain[i] = post (&mutex); wait (&mutex); state[i] = Think; captain[LEFT] = captain[RIGHT]; post (&mutex);?; if (i==1) i=i+1; else i=i+1; };

HF: void captain (int i) { if (state[i] == E; post (&mutex);?; if (state[i] != E) i=i+1; else i=i+1; };

State LF: ! = E; if (state[i] == E; post (&mutex);?; if (state[i] != E;

Wait & post
when SIGCHLD comes, the corresponding signal handling routine is registered in the parent with SIGCHLD. When SIGCHLD is issued, the parent will receive the signal. If it could be avoided, the problem of race exclusion is a requirement of race condition. If it could be avoided, the problem of race condition is solved.

critical section: A code segment that access shared objects simultaneously and can be designed for accessing more than one shared object.

File system calls -> **Entry and Exit Implementation**

fs
1. clean up the exit process's resource.
2. clean up the exit process's resource.
3. notify the parent with SIGCHLD when SIGCHLD is issued. The corresponding signal handling routine is invoked (remove the signal, destroy the child process condition variable). If it could be avoided, the problem of race exclusion is solved.

fs (entry)
1. deregister the signal handling routine in the kernel space.
2. for the parent 4. return the PID of child for the parent 4. return the PID of child for the parent 4. return the PID of child for the parent 4. return the PID of child for the parent 4. return the PID of child

fs (exit)
1. The process terminates abnormally.
2. The process terminates normally.
3. The process terminates normally.

When making a system call (fs) switches from user mode to kernel mode.

Time

User Time

CPU

idle

IO

SWI

SWI

SWI

System Call
Program state • PC, Register, Execution Flows, Stack, program state • User programs and user programs are isolated - different processes share the same physical address space • Program executes in an address space of the physical machine distinct from the memory space of the program's context of that system call = print, sort, move, free, fopen, fclose, write etc.

or **Fundamental OS Concept**
- Single unique execution context, fully describes the system state • PC, Register, Execution Flows, Stack, program state • User programs and user programs are isolated - different processes share the same physical address space • Program executes in an address space of the physical machine distinct from the memory space of the program's context of that system call = mdm, clm, open, chmod, read, write, etc.

Al mode operation/Protection
- Only the system has the listing of an address space and more threads of the system has the listing of hardware resources • The OS and the hardware resources are isolated - different processes spend most of time in CPU. (CPU-bound process) - spend most of time in I/O. (I/O-bound process) - spend most of time in I/O. (I/O-bound process) - spend most of time in I/O. (I/O-bound process)

Turnaround Time: Turnaround Time - task time Δt
Waiting Time: Waiting time for event finished Δt
DD sum: The responsiveness of the process is greater than Δt if the waiting time is less than the turnaround time.

Thread State
Saved: Heap, Global V, Program code, save PC, SP and register in current state
Switch from one CPU to the next:
in switch from one CPU to another, the context is saved and restored

• Allow system to enter deadlock and then recover. Prevent retain without letting off process, tell back action of deallocated threads if ($X \leq 0$) 'wake-up'; y enable-interrupt();

Four requirement for Deadlock
• Mutual Exclusion
• Hold and Wait
• Circular Wait
• No preemption
• Method of Handling Deadlock
• Allow system to enter deadlock and then recover. Ignore the deadlock when detects deadlock

What to do when detects deadlock
• Terminate thread, force it to give up resources. Prevent retain without letting off process, tell back action of deallocated threads if ($X \leq 0$) 'wake-up'; y enable-interrupt();

Thread: Infinite resource • No sharing of resource. Do not allow multiple threads to share the same resource. Make all threads request everything they need in the beginning.

Thread Creation: Ready thread yield / Scheduler suspend / Event occur / Thread lifecycle

Thread: Ready \rightarrow Running \rightarrow Zombie \rightarrow Death

Thread: Blocked \rightarrow Waiting

Thread: Blocked \rightarrow Ready \rightarrow Running

Thread: Blocked \rightarrow Ready \rightarrow Waiting

Thread: Blocked \rightarrow Ready \rightarrow Waiting