

What is OS

special layer of software that provides application software access to hardware resources:

Content abstraction of complex hardware device
Protected access to shared resources
Security and authentication
Communication among logical entities

OS includes

kernel manage all the physical devices
system call configure kernel or build things
driver handle the interaction
shell render a simple command line

OS do

• Provide abstractions to apps (File systems;

Process, threads; VM, containers)

• Manage resources (Memory, CPU, storage)

that is system call

• is a function call

• exposed by the kernel

• abstract away most low-level details

NOT system call

malloc, scanf, malloc, free, fopen, fclose

is system call

mkdir, rmdir, chown, chmod, open, fchdir,

chmod, getpid, pause, pipe, read, signal, stat,

sleep, wait, write

memory of a process



Four fundamental OS concepts

Thread • single unique execution context.

fully describes program state

• PC, Register, Execution Flags, Stack

Address Space • Programs execute in an

address space that is distinct from the

memory space of the physical machine.

Process • An instance of an executing

program is a process consisting of an

address space and 1 more thread of control

Dual mode operation/Protection • Only the

system has the ability to access certain

resources. • The OS and the hardware

are protected from user programs and

user programs are isolated from one

another by controlling the translation

from program virtual addresses to

machine physical addresses

CPU need structure to hold

PC, SP (Stack Pointer), Register

How switch from one CPU to the next

1. Save PC, SP and register in current state

2. Load PC, SP and register from new state block.

What triggers switch

Timer, voluntary yield, I/O

Types of Mode Transfer: Unprogrammed

System Call, Interrupt, Trap or Exception

Banker

Multithreaded Processes

[Avail] = [FreeResources]

Add all nodes Unfinished

done = true;

foreach node UF;

if ([Request] < [Avail])

|| [Max] - [Alloc] < [Avail]

remove node from UF

[Avail] = [Avail] + [Alloc]

done = false;



fork()

It creates the child process by cloning from the parent process including all user-space data: PCI CPU reg, Program code [File & Memory], Memory, Opened files [Kernel's internal]. However, DO NOT CLONE Return value of fork(), PID, Parent process, Running time, File locks.

exec()

The process is changing the code that is executing and never return to the original code. Will replace: Program Code, Memory (local v global v, dynamically allocated memory), Reg value (PC). However, the kernel-space info is preserved (PID, Process relationship). suspend the calling process to waiting state and return when • one of its child (running → terminated) • signal is received • return immediately if it has no child.

fork() inside the kernel

1. copy kernel space 2. update kernel space (PID, Running time, pointer to my parent / List of children) 3. Array of opened files → present

3. copy user space 4. return value

exec() inside the kernel

1. Clear local v & dynamically allocated memory, reset global v based on new code, change constants to the new program code

The kernel will also reset the register value (PC)

exit() inside the kernel

1. Clean up most of the allocated kernel space memory

2. Clean up the exit process's user space memory

3. notify the parent with SIGCHLD

State: zombie only has pid

wait() inside the kernel

1. register a signal handling routine 2. when SIGCHLD comes, the

corresponding signal handling routine is invoked (remove the signal)

destroy the child process in the kernel space) 3. deregister the

signal handling routine for the parent 4. return the PID of child.

exit() turns a process into a zombie when

1. The process calls exit()

2. The process return from main()

3. The process terminates abnormally

Multi-threaded Processes

- PCB point to multiple TCB
- switching thread within a block is simple
- switching thread across blocks requires changes to memory and I/O address tables

Multi-processing \equiv Multiple CPUs

Multi-programming \equiv Multiple Jobs / Process

Multi-threading \equiv Multiple threads per process

Multi-processing $\begin{cases} A \rightarrow B \\ B \rightarrow A \end{cases}$

Multi-programming $A \rightarrow B$

Threading Model

One-to-One, Many-to-One, Many-to-Many

IPC $\begin{cases} \text{FIFO} > \text{system} \\ \text{Pipe} < \text{semaphore} \rightarrow \text{user} \end{cases}$

Shared Objects

Message Passing

Signal

Race condition

Happen whenever shared object + multiple process + concurrently

The outcome of the computation depends on the execution sequences of the process

Mutual exclusion is a requirement

If it could be achieved, the problem of race condition would be gone.

A critical section

the code segment that access shared object.

can be designed for accessing more than one shared object

Entry and exit implementation

Natural Exclusion

Achieving Mutual Exclusion

- Spin-based lock
- Spin semaphore
- Semaphore

suffer from priority inversion

PC is pointing to the region (code) when in kernel
CPU switches from userspace to kernel space, reads the PID of the process from the kernel, it switches back to the userspace memory, continue running program code

User Time: CPU time spent on codes in userspace mem.

Sys Time: CPU time spent on codes in kernelspace mem.

Process Lifecycle

Forked \rightarrow Ready \rightleftharpoons Running \rightarrow Zombie

Block / Waiting same process: TCB, reg, PC, GP

Context Switch \rightarrow Save diff process: PCB, file-descriptor

the scheduler decides to schedule another process in the ready queue. Then the scheduler has to load the context of that process from the main memory to CPU

[save and restore reg • switch address space • cache, buffer cache & TLB misses]

PU-bound Process

pend most of time in CPU, user-time $>$ sys-time [AI]

IO-bound Process [real-time $>$ user-time + sys-time]

pend most of time in I/O, sys-time $>$ user-time [I/Os]

unround time Finish-time - Arrive-time

Waiting time Turnaround time - task-time

R: The responsiveness of the process is greater, not frozen

Thread State

shared State

Heap

Global variable

Code

Thread Control Block

Stack Info

Saved Reg

Thread Metadata

Stack



Spin semaphore

int turn; int interested [Z] = 1 FALSE, FALSE;
void lock (int process) {
int other; other = 1 - process;
turn = other; while (turn == other && interested[process] = TRUE);
void unlock (int process) { interested[process] = FALSE; }
→ A Low priority process L is inside the critical region, but a high priority process H gets the CPU and want to enter the critical region. But H cannot lock (since L has not unlock).

Producer-consumer

SO: semaphore mutex = 1; avail = N, fill = 0;

ES: void producer { int item; while (true) { item = produce_item();

wait (&avail); wait (&mutex); insert_item(item);

post (&mutex); post (&fill); } }

ES: void consumer { int item; while (true) { wait (&fill);

wait (&mutex); item = remove_item(); post (&mutex);

post (&avail); } }

Dining philosopher (LEFT ((i+N-1)%N) Right ((i+1)%N)

SO: int state[N]; semaphore mutex = 1, P[N] = 0

MF: void philosopher (int i) { think(); take (&i); eat(); put (&i); }

SE: void take (int i) { wait (&mutex); state[i] = Hungry;

captain(i); post (&mutex); wait (&P[i]); }

SE: void put (int i) { wait (&mutex); state[i] = Think;

captain(LEFT); captain(RIGHT); post (&mutex); }

HF: void captain (int i) { if (state[i] = H & state[L] = E & state

[R] != E) { state[i] = E; post (&P[i]); }

wait & post

void wait (semaphore *s) { disable_interrupt(); *s = *s - 1;

if (*s < 0) { enable_interrupt(); sleep(); disable_interrupt(); }

enable - interrupt(); }

void post (semaphore *s) { 1 disable_interrupt(); *s = *s + 1;

if (*s < 0) wake-up(); enable - interrupt(); }

Four requirement for DeadLock

- Mutual Exclusion
- Hold and Wait
- No preemption